

Policy comparison on test set

Policy	Accuracy (95% CI)	Mean chunks (95% CI)	Cost-adj. utility (95% CI)
FixedK(k=3)	100.0% (100–100%)	4.00 (4.00–4.00)	0.60 (0.60–0.60)
FixedK(k=5)	100.0% (100–100%)	6.00 (6.00–6.00)	0.40 (0.40–0.40)
Heuristic(0.8)	98.0% (92–100%)	3.64 (3.50–3.78)	0.62 (0.57–0.65)
Conservative QL ($\alpha=0.5$)	100.0% (100–100%)	4.10 (4.10–4.10)	0.59 (0.59–0.59)

Table 1. Policy comparison on test set. 50 episodes, seed 999. Cost-adj. utility = accuracy – 0.1 $\times$ mean chunks; 95% CIs from bootstrap (1000 samples).