



# Conservative Q-Learning Training Curves (Hyperparameter Sweep)


Total Loss




TD Loss (Bellman Error)



Conservative Penalty



Mean Q-Value



—  $\alpha = 0.5, lr = 1e-3$  (best)    —  $\alpha = 1.0, lr = 3e-4$     —  $\alpha = 0.1, lr = 3e-4$