

Conservative Q-Learning training curves (hyperparameter sweep)

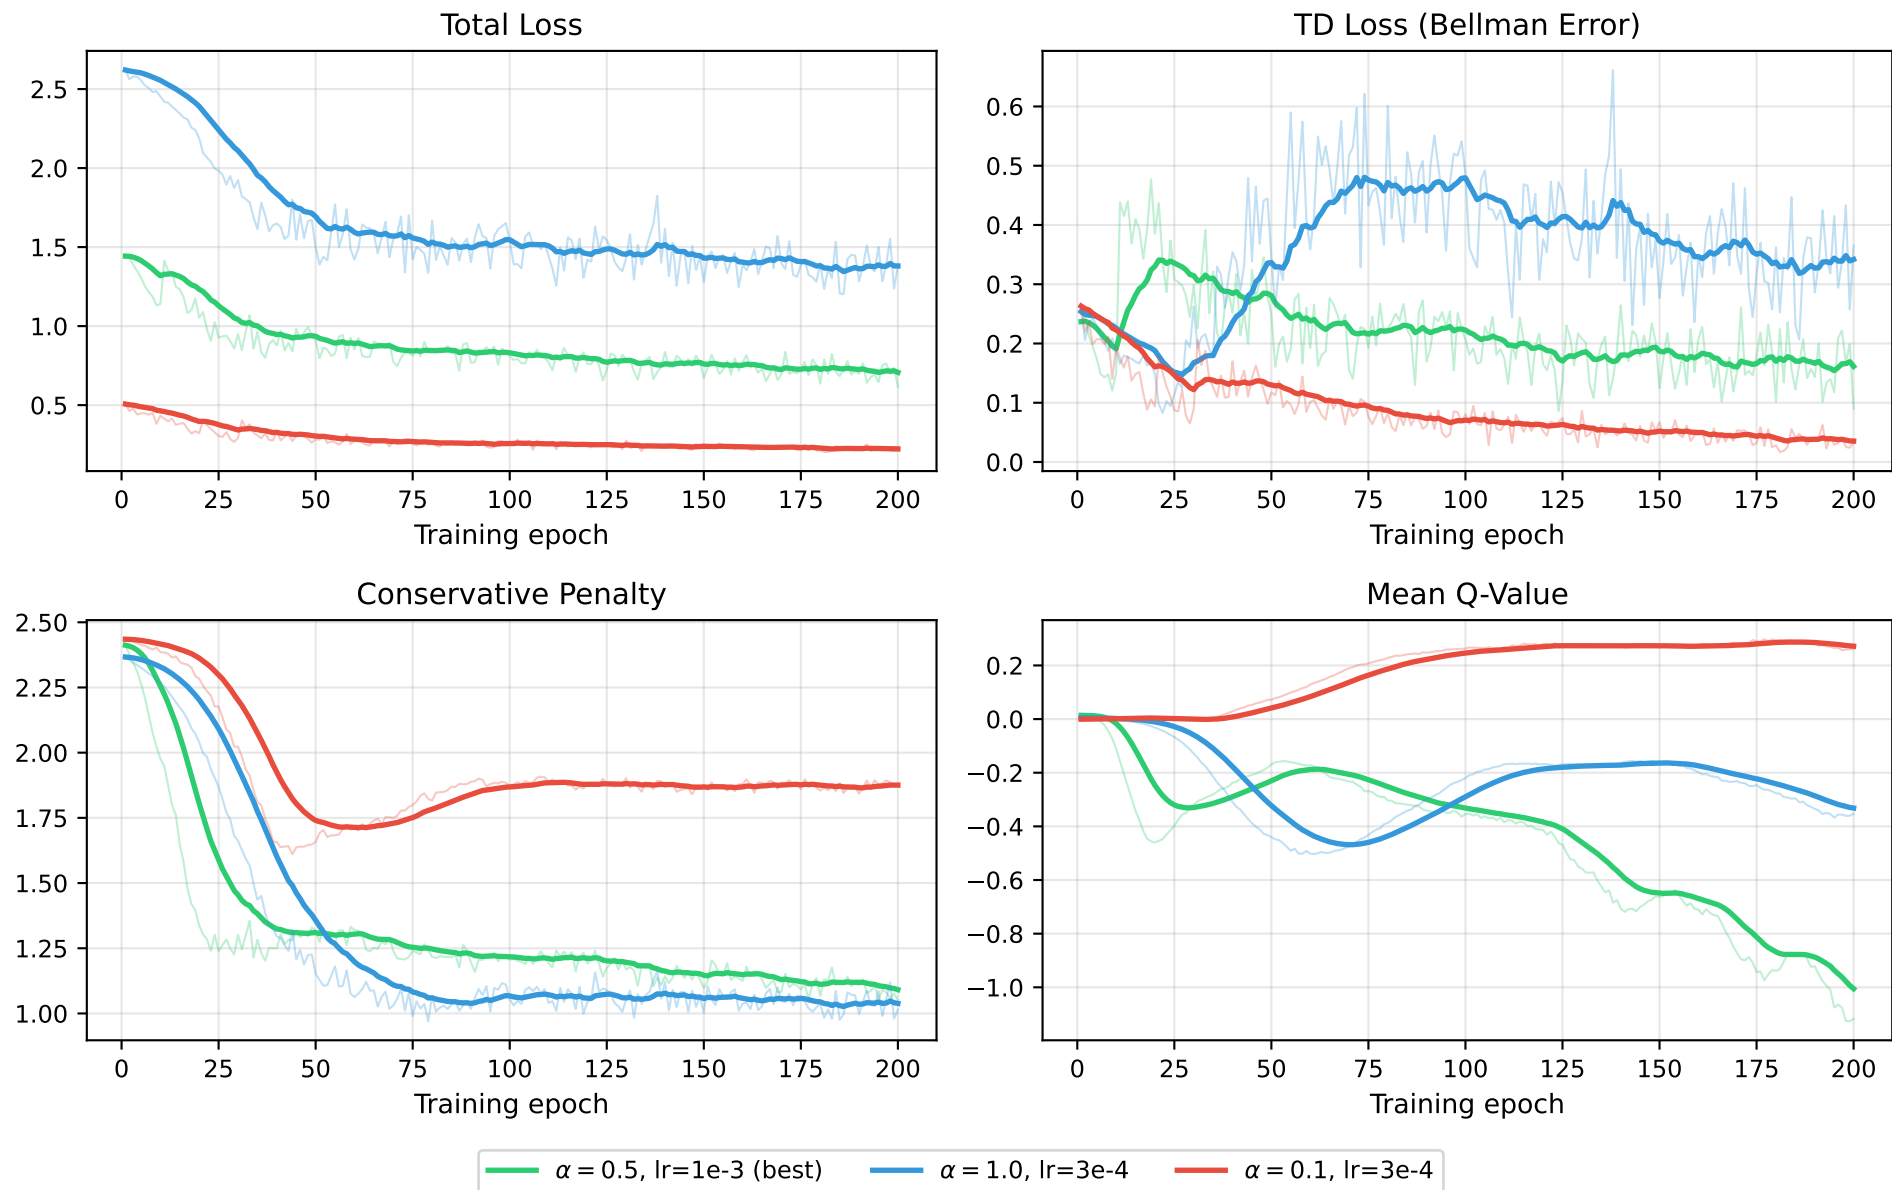


Figure 2. Conservative Q-Learning training curves (hyperparameter sweep). Solid lines: exponential moving average ($\alpha = 0.9$). Best config ($\alpha = 0.5, lr = 1e-3$) balances accuracy and efficiency.