



Training Curves: CQL Hyperparameter Sweep (α , lr)


Total Loss




TD Loss (Bellman Error)



Conservative Penalty



Mean Q-Value



— CQL $\alpha = 0.5, lr = 1e-3$ (best) — CQL $\alpha = 1.0, lr = 3e-4$ — CQL $\alpha = 0.1, lr = 3e-4$