



Training Curves: CQL Hyperparameter Sweep (α , lr)


Total Loss




TD Loss (Bellman Error)



Conservative Penalty



Mean Q-Value



— CQL $\alpha = 0.5$, lr=1e-3 (best) — CQL $\alpha = 0.1$, lr=3e-4 — CQL (2k corpus)
— CQL $\alpha = 1.0$, lr=3e-4